

Runtime Verification of Smart Contracts

On the Ethereum network

Lars Stegeman [s1346466]
l.stegeman@student.utwente.nl

August 28, 2018

Contents

1	Abstract	1
2	Introduction	1
3	Solidity	1
4	Tool Overview	1
5	Annotation Language	2
5.1	ANTLR	2
5.2	Grammar definition	2
6	Validation	4
6.1	Design	4
6.2	Implementation	4
7	Generation	4
8	Case study	4
9	Future work	4
10	Related Work	5

1 Abstract

2 Introduction

3 Solidity

Section explains basics of solidity in order to understand the rest of the paper. Same section as in ResearchTopics

4 Tool Overview

- Tool works in two phases validation and generation.
- Validation typechecks and checks if annotations are well formed according to grammar.
- Generation generates original solidity code with extra added code to check annotations at runtime.
- Approach uses a Solidity grammar that can be updated easily for future updates.

- Output of validation phase can be used for other verification tools (result is a parse tree).
- Annotations can be defined as invariants and pre or post conditions for functions.
- Annotations use JML like syntax.
- Tool can be used both during development (as extra test cases) or on the actual live blockchain (this probably costs a lot of extra gas = ether).

5 Annotation Language

The first step towards implementing the tool is to define an annotation syntax, and formally write this down using a grammar.

5.1 ANTLR

5.2 Grammar definition

There already exists a grammar for the complete Solidity language. This grammar is written down using the language that is used by ANTLR tool. ANTLR is a parser generator which helps automate the building of a compiler. ANTLR has the capabilities to extend certain grammars. This is done by inheritance over the original grammar. This principle is explained in detail here (find ref). We will use this principle to extend the grammar of Solidity to recognize the special annotations that will later be used in the tool.

The annotations have certain requirements that can be summarized the following way. Later each requirement is discussed in detail.

- Annotations can be specified at the top level.
- Annotations should be able to reference all variables used in the contract.
- Basic math operations can be used within annotations.
- Annotations can not have side effects.
- The type should be boolean at the highest level (that way they can be verified).
- There are two types of annotations invariants and pre- or postconditions to a function.
- Global variables (`msg` and `address`) can be used in annotations.

The annotation syntax is heavily inspired from the JML annotation syntax. But has a lot less built in keywords since the setting is easier and the tool is less complex. The original grammar is extended in such a way that annotations can only be defined on the top level. The relevant parts of the original Solidity grammar can be seen in the snippet below. This does not include the full grammar specification but only the parts that are relevant for the annotation syntax.

```
grammar Solidity;

sourceUnit
    : (pragmaDirective | importDirective | contractDefinition)* EOF ;

contractDefinition
    : ( 'contract' | 'interface' | 'library' ) identifier
      ( 'is' inheritanceSpecifier (',' inheritanceSpecifier)* )?
      '{' contractPart* '}' ;

contractPart
    : stateVariableDeclaration
    | usingForDeclaration
    | structDefinition
    | constructorDefinition
    | modifierDefinition
```

```

| functionDefinition
| eventDefinition
| enumDefinition ;

```

In the original grammar the definition of `contractPart` is what defines the declaration of variables and the definitions for structs and functions. This is where the extra annotations have to be added to the grammar. The snippet below shows the basic definition of an annotation. This is not the complete grammar some of the tokens are omitted from this snippet, since they are not required to understand the grammar definition.

```

grammar SolidityAnnotated;
import Solidity;

contractPart
: stateVariableDeclaration
| usingForDeclaration
| structDefinition
| constructorDefinition
| modifierDefinition
| functionDefinition
| eventDefinition
| enumDefinition
| annotationDefinition ;

annotationDefinition
: AnnotationStart AnnotationKind annotationExpression ;

annotationExpression
: '(' annotationExpression ')'
| annotationExpression compareOp annotationExpression
| annotationExpression booleanOp annotationExpression
| annotationExpression integerOpBoolean annotationExpression
| annotationExpression integerOpInteger annotationExpression
| '!' annotationExpression
| ('\\forall' | '\\exists') '(' identifier elementaryTypeName ':' annotationExp
| primaryAnnotationExpression ;

primaryAnnotationExpression
: primaryExpression
| primaryAnnotationExpression '.' identifier
| primaryAnnotationExpression '[' primaryExpression ']'
| '\\old' '(' primaryAnnotationExpression ')' ;

AnnotationStart
: '//@' ;

AnnotationKind
: 'inv' | 'pre' | 'post' ;

LINE_COMMENT
: '//' ~[@] ~[\\r\\n]* -> channel(HIDDEN) ;

```

An `AnnotationDefinition` is composed of multiple components. It has a `AnnotationStart`, `AnnotationKind` and `annotationExpression` component. The `AnnotationStart` token is used to signal that an annotation definition is coming next. This is defined as `//@` making it a line comment to other solidity compilers. This makes annotated solidity code still compilable by normal compilers. For the grammar to accept this notation the `LINE_COMMENT` token has to be adjusted to not accept `@` as a second character. Otherwise all annotation comments would be recognized as a `LINE_COMMENT` making it unusable.

There are two types of annotations they are defined by the token `AnnotationKind`. They can

either be a invariant or a pre- or post-condition of a function.

Each annotation has an expression which has to be evaluated called `annotationExpression`. The expression parser rules are separated between `annotationExpression` and `primaryAnnotationExpression`. This is needed to keep the hierarchy in parsing and prevent using primary definitions with complex expressions. For example using the keyword `'\old'` before parenthesis.

The annotation expressions use a different parser rules than the expression rules that are used within the original Solidity grammar. The `annotationExpression` does not allow syntax like `expression + '++'` and to distinguish these a new parser rule was introduced for annotations only.

6 Validation

6.1 Design

Annotations have to be validated on certain aspects for them to be correct and usable. These aspects have to be verified first for the annotations to be useful in the next generation phase. The parser ensures annotations are syntactically correct however there are more properties that have to be checked. In short the validation step within the program has to do the following things:

- Collect all variables within the contract.
- Typecheck annotations.

The validation phase will consist of two programs that walk the complete parse tree. The first walk will collect all the variables and defined structures and store these in an information object. The second walk will type check each annotation individually. During this type checking the type of each identifier is looked up using the collected information from the first walk.

6.2 Implementation

7 Generation

- Explain generation of functions for annotations. Each annotation will be transformed to a function, the function takes in arguments that are also generated automatically.
- Explain how functions are calling the code for checking annotations. Original Function will be transformed to a private function and renamed to `'functionName'+ _body`. New function calls this function with the extra annotations.
- Implementation of generation is using a `TokenStreamRewriter` (with whitespace) and an abstract parse tree (only tokens that are in grammar). Tokens in both structures know the position, this way the original solidity code can be printed and extra code can be added through traversing the parse tree.

8 Case study

- Take contract that has vulnerability and add annotations
- Show parts of generated code that will throw an event.
- Show difference in gas usage?
- Show how forall annotations are handled (example contract of minimal token).

9 Future work

- The tool cannot handle contract inheritance.
- The tool does not have a mechanism for exposing the parse tree for other programs after validation.

10 Related Work

- ContractLarva: runtime monitoring based on defined states and transitions.
-